

# CNN\_3-BNDO

April 26, 2017

## 1 CS109B Project Group 26 - Deep Learning

## Main Specifications - Ver5

**\*\* Data preprocessing:\*\***

- Channel rearrangement
- 14 labels multi label classification
- data augmentation by shift and zoom for 8000 training samples and 2000 test samples
- Centered features

**\*\* Main Architecture \*\***

- Multi layer CNN with Batch Normalization Layers and Dropout :

- Conv2D - Relu depth 32 and 7x7 kernel - He uniform initialization - 30 % Dropout - Batch Norm - MaxPool 2x2 - Conv2D - Relu depth 32 and 5x5 kernel - He uniform initialization - 30 % Dropout - Batch Norm - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - 30 % Dropout - Batch Norm - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - 30 % Dropout - Batch Norm - FC 128 , Relu , He uniform initialization - 50 % Dropout - Batch Norm - FC 64 , Relu , He uniform initialization - 50 % Dropout - Batch Norm

- SGD optimizer with 1e-4 learning rate and 0.99 momentum - Binary cross entropy loss function - Batch size 128 - Training convergence after 100 epochs

### 1.0.1 Import Modules

```
In [1]: import requests
import json
import time
import itertools
import wget
import os
import pickle
import numpy as np

import random
import matplotlib
import matplotlib.pyplot as plt
```

```

%matplotlib inline

import seaborn as sns
from sklearn.cluster.bicluster import SpectralCoclustering
from sklearn.metrics import precision_recall_curve
import scipy

sns.set_style('white')
import tensorflow as tf
import pandas as pd
import keras
from keras.optimizers import SGD, Adam
from keras.models import Sequential
from keras.layers import Dense,Dropout,Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
import keras.initializers as init
from keras.preprocessing.image import ImageDataGenerator
from keras import backend as K
from keras.models import load_model

```

Using TensorFlow backend.

```
In [2]: x_train_dict = pickle.load(open('training_num.pik' , 'rb'))
```

```
In [3]: train_split = 8000
```

```
In [4]: x_train_raw = x_train_dict['images']
```

```

x_train = np.array(x_train_raw)[:train_split,: ,: , :]
x_test  = np.array(x_train_raw)[train_split,: ,: , :]

print x_train.shape

```

```
(8000, 128, 85, 3)
```

```
In [5]: img_rows = x_train.shape[1]
```

```
img_cols = x_train.shape[2]
```

```

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 3, img_rows, img_cols)
    input_shape = (3, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 3)
    input_shape = (img_rows, img_cols, 3)

```

```
In [6]: x_train = x_train.astype('float32')
        x_test  = x_test.astype('float32')

        x_train /= 255.0
        x_test  /= 255.0

        print 'x_train shape:', x_train.shape
        print x_train.shape[0], 'train samples'

        print 'x_test shape:', x_test.shape
        print x_test.shape[0], 'test samples'
```

```
x_train shape: (8000, 128, 85, 3)
8000 train samples
x_test shape: (1988, 128, 85, 3)
1988 test samples
```

```
In [7]: y_raw = pd.read_csv('Genres_labels_All_cleaned.csv')

        y_train = y_raw.iloc[:, 1:-1].values[:train_split, :]
        y_test  = y_raw.iloc[:, 1:-1].values[train_split:, :]

        num_classes = y_train.shape[1]

        print 'number of classes:  ', num_classes
```

```
number of classes:  14
```

```
In [8]: datagen = ImageDataGenerator(
        featurewise_center=True,
        featurewise_std_normalization=True,
        width_shift_range=0.2,
        height_shift_range=0.2,
        zoom_range = 0.5,
        fill_mode = 'wrap')

        datagen.fit(x_train)

        datagen.fit(x_test)
```

```
In [34]: # create an empty network model
        modelb = Sequential()

        # --- input layer ---
        modelb.add(Conv2D(32, kernel_size=(7, 7), activation='relu', input_shape=input_shape
                           kernel_initializer = init.he_normal(109)))
```

```

# -----Dropout -----
modelb.add(Dropout(0.3))

# -----Batch Normalization -----
modelb.add(BatchNormalization())

# --- max pool ---
modelb.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# ----- Conv Layer -----
modelb.add(Conv2D(32, kernel_size=(5, 5), activation='relu',
                  kernel_initializer = init.he_normal(109)))

# -----Dropout -----
modelb.add(Dropout(0.3))

# -----Batch Normalization -----
modelb.add(BatchNormalization())

# --- max pool ---
modelb.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# --- Conv layer ---
modelb.add(Conv2D(64, kernel_size=(3, 3), activation='relu',
                  kernel_initializer = init.he_normal(109)))

# -----Dropout -----
modelb.add(Dropout(0.3))

# -----Batch Normalization -----
modelb.add(BatchNormalization())

# --- max pool ---
modelb.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# --- Conv layer ---
modelb.add(Conv2D(64, kernel_size=(3, 3), activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# -----Dropout -----

```

```

modelb.add(Dropout(0.3))

# -----Batch Normalization -----
modelb.add(BatchNormalization())

# --- max pool ---
modelb.add(MaxPooling2D(pool_size=(2, 2)))

#-----
# flatten for fully connected classification layer
modelb.add(Flatten())

# --- fully connected layer ---
modelb.add(Dense(128, activation='relu',
                 kernel_initializer = init.he_normal(109)))

# -----Dropout -----
modelb.add(Dropout(0.5))

# -----Batch Normalization -----
modelb.add(BatchNormalization())

# -----
# --- fully connected layer ---
modelb.add(Dense(64, activation='relu' ,
                 kernel_initializer = init.he_normal(109)))

# -----Dropout -----
modelb.add(Dropout(0.5))

# -----Batch Normalization -----
modelb.add(BatchNormalization())

# -----
# --- classification ---
modelb.add(Dense(num_classes, activation='sigmoid'))

# prints out a summary of the model architecture
modelb.summary()

```

| Layer (type)       | Output Shape        | Param # |
|--------------------|---------------------|---------|
| conv2d_21 (Conv2D) | (None, 122, 79, 32) | 4736    |

|  |                     |       |
|--|---------------------|-------|
| dropout_31 (Dropout)                         | (None, 122, 79, 32) | 0     |
| batch_normalization_31 (Batch Normalization) | (None, 122, 79, 32) | 128   |
| max_pooling2d_21 (MaxPooling2D)              | (None, 61, 39, 32)  | 0     |
| conv2d_22 (Conv2D)                           | (None, 57, 35, 32)  | 25632 |
| dropout_32 (Dropout)                         | (None, 57, 35, 32)  | 0     |
| batch_normalization_32 (Batch Normalization) | (None, 57, 35, 32)  | 128   |
| max_pooling2d_22 (MaxPooling2D)              | (None, 28, 17, 32)  | 0     |
| conv2d_23 (Conv2D)                           | (None, 26, 15, 64)  | 18496 |
| dropout_33 (Dropout)                         | (None, 26, 15, 64)  | 0     |
| batch_normalization_33 (Batch Normalization) | (None, 26, 15, 64)  | 256   |
| max_pooling2d_23 (MaxPooling2D)              | (None, 13, 7, 64)   | 0     |
| conv2d_24 (Conv2D)                           | (None, 11, 5, 64)   | 36928 |
| dropout_34 (Dropout)                         | (None, 11, 5, 64)   | 0     |
| batch_normalization_34 (Batch Normalization) | (None, 11, 5, 64)   | 256   |
| max_pooling2d_24 (MaxPooling2D)              | (None, 5, 2, 64)    | 0     |
| flatten_6 (Flatten)                          | (None, 640)         | 0     |
| dense_16 (Dense)                             | (None, 128)         | 82048 |
| dropout_35 (Dropout)                         | (None, 128)         | 0     |
| batch_normalization_35 (Batch Normalization) | (None, 128)         | 512   |
| dense_17 (Dense)                             | (None, 64)          | 8256  |
| dropout_36 (Dropout)                         | (None, 64)          | 0     |
| batch_normalization_36 (Batch Normalization) | (None, 64)          | 256   |
| dense_18 (Dense)                             | (None, 14)          | 910   |
| =====  |                     |       |
| Total params: 178,542.0                      |                     |       |
| Trainable params: 177,774.0                  |                     |       |

# CNN\_3-BNDOReg

April 26, 2017

## 1 CS109B Project Group 26 - Deep Learning

## Main Specifications - Ver6

**\*\* Data preprocessing:\*\***

- Channel rearrangement
- 14 labels multi label classification
- data augmentation by shift and zoom for 8000 training samples and 2000 test samples
- Centered features

**\*\* Main Architecture \*\***

- Multi layer CNN with Batch Normalization Layers and Dropout and 0.1 L2 Kernel Regularization :

- Conv2D - Relu depth 32 and 7x7 kernel - He uniform initialization - 0.1 L2 Kernel Regularization - 30 % Dropout - Batch Norm - MaxPool 2x2 - Conv2D - Relu depth 32 and 5x5 kernel - He uniform initialization - 0.1 L2 Kernel Regularization - 30 % Dropout - Batch Norm - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - 0.1 L2 Kernel Regularization - 30 % Dropout - Batch Norm - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - 0.1 L2 Kernel Regularization - 30 % Dropout - Batch Norm - FC 128 , Relu , He uniform initialization - 0.1 L2 Kernel Regularization - 50 % Dropout - Batch Norm - FC 64 , Relu , He uniform initialization - 0.1 L2 Kernel Regularization - 50 % Dropout - Batch Norm

- SGD optimizer with 1e-4 learning rate and 0.99 momentum - Binary cross entropy loss function - Batch size 128 - Training convergence after 100 epochs

### 1.0.1 Import Modules

```
In [1]: import requests
import json
import time
import itertools
import wget
import os
import pickle
import numpy as np
```

```

import random
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

import seaborn as sns
from sklearn.cluster.bicluster import SpectralCoclustering
from sklearn.metrics import precision_recall_curve
import scipy

sns.set_style('white')
import tensorflow as tf
import pandas as pd
import keras
from keras import regularizers as reg
from keras.optimizers import SGD, Adam
from keras.models import Sequential
from keras.layers import Dense,Dropout,Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
import keras.initializers as init
from keras.preprocessing.image import ImageDataGenerator
from keras import backend as K
from keras.models import load_model

```

Using TensorFlow backend.

### Open the Preprocessed Poster Data

```
In [2]: x_train_dict = pickle.load(open('training_num.pik' , 'rb'))
```

### Specify the training/test split

```
In [3]: train_split = 8000
```

```
In [4]: x_train_raw = x_train_dict['images']
```

```

x_train = np.array(x_train_raw)[:train_split,: ,: , :]
x_test  = np.array(x_train_raw)[train_split:,: ,: , :]

```

```
print x_train.shape
```

```
(8000, 128, 85, 3)
```

```
In [5]: img_rows = x_train.shape[1]
```

```
img_cols = x_train.shape[2]
```



```

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 3, img_rows, img_cols)
    input_shape = (3, img_rows, img_cols)

else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 3)
    input_shape = (img_rows, img_cols, 3)

In [6]: x_train = x_train.astype('float32')
        x_test  = x_test.astype('float32')

        x_train /= 255.0
        x_test  /= 255.0

        print 'x_train shape:', x_train.shape
        print x_train.shape[0], 'train samples'

        print 'x_test shape:', x_test.shape
        print x_test.shape[0], 'test samples'

x_train shape: (8000, 128, 85, 3)
8000 train samples
x_test shape: (1988, 128, 85, 3)
1988 test samples

In [7]: y_raw = pd.read_csv('Genres_labels_All_cleaned.csv')

        y_train = y_raw.iloc[:, 1:-1].values[:train_split, :]
        y_test  = y_raw.iloc[:, 1:-1].values[train_split:, :]

        num_classes = y_train.shape[1]

        print 'number of classes:  ', num_classes

number of classes:  14

In [8]: datagen = ImageDataGenerator(
        featurewise_center=True,
        featurewise_std_normalization=True,
        width_shift_range=0.2,
        height_shift_range=0.2,
        zoom_range = 0.5,
        fill_mode = 'wrap')

        datagen.fit(x_train)

        datagen.fit(x_test)

```

```

In [9]: # Specify regularization parameter
        reg_par = 0.1

In [10]: # create an empty network model
        modelc = Sequential()

        # --- input layer ---
        modelc.add(Conv2D(32, kernel_size=(7, 7), activation='relu', input_shape=input_shape,
                           kernel_initializer = init.he_normal(109),
                           kernel_regularizer = reg.l2(reg_par)))

        # -----Dropout -----
        modelc.add(Dropout(0.3))

        # -----Batch Normalization -----
        modelc.add(BatchNormalization())

        # --- max pool ---
        modelc.add(MaxPooling2D(pool_size=(2, 2)))

        # -----
        # -----
        # --- Conv Layer ---
        modelc.add(Conv2D(32, kernel_size=(5, 5), activation='relu',
                           kernel_initializer = init.he_normal(109),
                           kernel_regularizer = reg.l2(reg_par)))

        # -----Dropout -----
        modelc.add(Dropout(0.3))

        # -----Batch Normalization -----
        modelc.add(BatchNormalization())

        # --- max pool ---
        modelc.add(MaxPooling2D(pool_size=(2, 2)))

        # -----
        # -----
        # --- Conv layer ---
        modelc.add(Conv2D(64, kernel_size=(3, 3), activation='relu',
                           kernel_initializer = init.he_normal(109),
                           kernel_regularizer = reg.l2(reg_par)))

        # -----Dropout -----
        modelc.add(Dropout(0.3))

```

```

# -----Batch Normalization -----
modelc.add(BatchNormalization())

# --- max pool ---
modelc.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# --- Conv layer ---
modelc.add(Conv2D(64, kernel_size=(3, 3), activation='relu' ,
                  kernel_initializer = init.he_normal(109),
                  kernel_regularizer = reg.l2(reg_par)))

# -----Dropout -----
modelc.add(Dropout(0.3))

# -----Batch Normalization -----
modelc.add(BatchNormalization())

# --- max pool ---
modelc.add(MaxPooling2D(pool_size=(2, 2)))

#-----
# flatten for fully connected classification layer
modelc.add(Flatten())

# --- fully connected layer ---
modelc.add(Dense(128, activation='relu',
                 kernel_initializer = init.he_normal(109),
                 kernel_regularizer = reg.l2(reg_par)))

# -----Dropout -----
modelc.add(Dropout(0.5))

# -----Batch Normalization -----
modelc.add(BatchNormalization())

# -----
# --- fully connected layer ---
modelc.add(Dense(64, activation='relu' ,
                 kernel_initializer = init.he_normal(109),
                 kernel_regularizer = reg.l2(reg_par)))

# -----Dropout -----
modelc.add(Dropout(0.5))

```

```

# -----Batch Normalization -----
modelc.add(BatchNormalization())

# -----

# --- classification ---
modelc.add(Dense(num_classes, activation='sigmoid'))

# prints out a summary of the model architecture
modelc.summary()

```

| Layer (type)                                | Output Shape        | Param # |
|---|---------------------|---------|
| conv2d_1 (Conv2D)                           | (None, 122, 79, 32) | 4736    |
| dropout_1 (Dropout)                         | (None, 122, 79, 32) | 0       |
| batch_normalization_1 (Batch Normalization) | (None, 122, 79, 32) | 128     |
| max_pooling2d_1 (MaxPooling2D)              | (None, 61, 39, 32)  | 0       |
| conv2d_2 (Conv2D)                           | (None, 57, 35, 32)  | 25632   |
| dropout_2 (Dropout)                         | (None, 57, 35, 32)  | 0       |
| batch_normalization_2 (Batch Normalization) | (None, 57, 35, 32)  | 128     |
| max_pooling2d_2 (MaxPooling2D)              | (None, 28, 17, 32)  | 0       |
| conv2d_3 (Conv2D)                           | (None, 26, 15, 64)  | 18496   |
| dropout_3 (Dropout)                         | (None, 26, 15, 64)  | 0       |
| batch_normalization_3 (Batch Normalization) | (None, 26, 15, 64)  | 256     |
| max_pooling2d_3 (MaxPooling2D)              | (None, 13, 7, 64)   | 0       |
| conv2d_4 (Conv2D)                           | (None, 11, 5, 64)   | 36928   |
| dropout_4 (Dropout)                         | (None, 11, 5, 64)   | 0       |
| batch_normalization_4 (Batch Normalization) | (None, 11, 5, 64)   | 256     |
| max_pooling2d_4 (MaxPooling2D)              | (None, 5, 2, 64)    | 0       |
| flatten_1 (Flatten)                         | (None, 640)         | 0       |

# CNN\_3

April 26, 2017

## 1 CS109B Project Group 26 - Deep Learning

## Main Specifications - Ver2

**\*\* Data preprocessing:\*\***

- Channel rearrangement
- 14 labels multi label classification
- data augmentation by shift and zoom for 8000 training samples and 2000 test samples
- Centered features

**\*\* Main Architecture \*\***

- Multi layer CNN :

- Conv2D - Relu depth 32 and 7x7 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 32 and 5x5 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - FC 128 , Relu , He uniform initialization - FC 64 , Relu , He uniform initialization

- SGD optimizer with 1e-6 learning rate and 0.99 momentum - Binary cross entropy loss function - Batch size 64 - Training convergence after 30 epochs

### 1.0.1 Import Modules

```
In [1]: import requests
import json
import time
import itertools
import wget
import os
import pickle
import numpy as np

import random
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```

import seaborn as sns
from sklearn.cluster.bicluster import SpectralCoclustering
from sklearn.metrics import precision_recall_curve
import scipy

sns.set_style('white')
import tensorflow as tf
import pandas as pd
import keras
from keras.optimizers import SGD, Adam
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
import keras.initializers as init
from keras.preprocessing.image import ImageDataGenerator
from keras import backend as K
from keras.models import load_model

```

Using TensorFlow backend.

```
In [2]: x_train_dict = pickle.load(open('training_num.pik' , 'rb'))
```

```
In [3]: train_split = 8000
```

```
In [4]: x_train_raw = x_train_dict['images']
```

```

x_train = np.array(x_train_raw)[:train_split,: , : , :]
x_test  = np.array(x_train_raw)[train_split:,: , : , :]

```

```
print x_train.shape
```

```
(8000, 128, 85, 3)
```

```
In [5]: img_rows = x_train.shape[1]
```

```
img_cols = x_train.shape[2]
```

```

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 3, img_rows, img_cols)
    input_shape = (3, img_rows, img_cols)

```

```

else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 3)
    input_shape = (img_rows, img_cols, 3)

```

```

In [6]: x_train = x_train.astype('float32')
        x_test  = x_test .astype('float32')

```

```

x_train /= 255.0
x_test  /= 255.0

print 'x_train shape:', x_train.shape
print x_train.shape[0], 'train samples'

print 'x_test shape:', x_test.shape
print x_test.shape[0], 'test samples'

x_train shape: (8000, 128, 85, 3)
8000 train samples
x_test shape: (1988, 128, 85, 3)
1988 test samples

In [7]: y_raw = pd.read_csv('Genres_labels_All_cleaned.csv')

y_train = y_raw.iloc[:, 1:-1].values[:train_split, :]
y_test  = y_raw.iloc[:, 1:-1].values[train_split:, :]

num_classes = y_train.shape[1]

In [30]: datagen = ImageDataGenerator(
        featurewise_center=True,
        featurewise_std_normalization=True,
        width_shift_range=0.2,
        height_shift_range=0.2,
        zoom_range = 0.5,
        fill_mode = 'wrap')

datagen.fit(x_train)

datagen.fit(x_test)

In [31]: # create an empty network model
model2 = Sequential()

# --- input layer ---
model2.add(Conv2D(32, kernel_size=(7, 7), activation='relu', input_shape=input_shape,
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model2.add(MaxPooling2D(pool_size=(2, 2)))

# ---- Conv Layer ---
model2.add(Conv2D(32, kernel_size=(5, 5), activation='relu',
                  kernel_initializer = init.he_normal(109)))

```

```

# --- max pool ---
model2.add(MaxPooling2D(pool_size=(2, 2)))

# --- Conv layer ---
model2.add(Conv2D(64, kernel_size=(3, 3), activation='relu',
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model2.add(MaxPooling2D(pool_size=(2, 2)))

# --- Conv layer ---
model2.add(Conv2D(64, kernel_size=(3, 3), activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model2.add(MaxPooling2D(pool_size=(2, 2)))

# flatten for fully connected classification layer
model2.add(Flatten())

# --- fully connected layer ---
model2.add(Dense(128, activation='relu',
                  kernel_initializer = init.he_normal(109)))

# --- fully connected layer ---
model2.add(Dense(64, activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# --- classification ---
model2.add(Dense(num_classes, activation='sigmoid'))

# prints out a summary of the model architecture
model2.summary()

```

| Layer (type)                  | Output Shape        | Param # |
|-------------------------------|---------------------|---------|
| conv2d_13 (Conv2D)            | (None, 122, 79, 32) | 4736    |
| max_pooling2d_13 (MaxPooling) | (None, 61, 39, 32)  | 0       |
| conv2d_14 (Conv2D)            | (None, 57, 35, 32)  | 25632   |
| max_pooling2d_14 (MaxPooling) | (None, 28, 17, 32)  | 0       |
| conv2d_15 (Conv2D)            | (None, 26, 15, 64)  | 18496   |



# CNN\_4

April 26, 2017

## 1 CS109B Project Group 26 - Deep Learning

## Main Specifications - Ver3

**\*\* Data preprocessing:\*\***

- Channel rearrangement
- 14 labels multi label classification
- Data augmentation by shift and zoom for 8000 training sample and 2000 test sample
- Centered features

**\*\* Main Architecture \*\***

- Multi layer CNN :

- Conv2D - Relu depth 32 and 7x7 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 32 and 5x5 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - FC 128 , Relu , He uniform initialization - FC 64 , Relu , He uniform initialization

- Adam optimizer with 1e-6 learning rate - Binary cross entropy loss function - Batch size 64 - Training convergence after 30 epochs

### 1.0.1 Import Modules

```
In [1]: import requests
import json
import time
import itertools
import wget
import os
import pickle
import numpy as np

import random
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```

import seaborn as sns
from sklearn.cluster.bicluster import SpectralCoclustering
from sklearn.metrics import precision_recall_curve
import scipy

sns.set_style('white')
import tensorflow as tf
import pandas as pd
import keras
from keras.optimizers import SGD, Adam
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
import keras.initializers as init
from keras.preprocessing.image import ImageDataGenerator
from keras import backend as K
from keras.models import load_model

```

Using TensorFlow backend.

```
In [2]: x_train_dict = pickle.load(open('training_num.pik' , 'rb'))
```

```
In [3]: train_split = 8000
```

```
In [4]: x_train_raw = x_train_dict['images']
```

```

x_train = np.array(x_train_raw)[:train_split,: , : , :]
x_test  = np.array(x_train_raw)[train_split:,: , : , :]

```

```
print x_train.shape
```

```
(8000, 128, 85, 3)
```

```
In [5]: img_rows = x_train.shape[1]
```

```
img_cols = x_train.shape[2]
```

```

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 3, img_rows, img_cols)
    input_shape = (3, img_rows, img_cols)

```

```

else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 3)
    input_shape = (img_rows, img_cols, 3)

```

```

In [6]: x_train = x_train.astype('float32')
        x_test  = x_test .astype('float32')

```

```

x_train /= 255.0
x_test  /= 255.0

print 'x_train shape:', x_train.shape
print x_train.shape[0], 'train samples'

print 'x_test shape:', x_test.shape
print x_test.shape[0], 'test samples'

x_train shape: (8000, 128, 85, 3)
8000 train samples
x_test shape: (1988, 128, 85, 3)
1988 test samples

In [7]: y_raw = pd.read_csv('Genres_labels_All_cleaned.csv')

y_train = y_raw.iloc[:, 1:-1].values[:train_split, :]
y_test  = y_raw.iloc[:, 1:-1].values[train_split:, :]

num_classes = y_train.shape[1]

In [8]: datagen = ImageDataGenerator(
        featurewise_center=True,
        featurewise_std_normalization=True,
        width_shift_range=0.2,
        height_shift_range=0.2,
        zoom_range = 0.5,
        fill_mode = 'wrap')

datagen.fit(x_train)

datagen.fit(x_test)

In [10]: # create an empty network model
model3 = Sequential()

# --- input layer ---
model3.add(Conv2D(32, kernel_size=(7, 7), activation='relu', input_shape=input_shape,
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model3.add(MaxPooling2D(pool_size=(2, 2)))

# ---- Conv Layer ---
model3.add(Conv2D(32, kernel_size=(5, 5), activation='relu',
                  kernel_initializer = init.he_normal(109)))

```

```

# --- max pool ---
model3.add(MaxPooling2D(pool_size=(2, 2)))

# --- Conv layer ---
model3.add(Conv2D(64, kernel_size=(3, 3), activation='relu',
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model3.add(MaxPooling2D(pool_size=(2, 2)))

# --- Conv layer ---
model3.add(Conv2D(64, kernel_size=(3, 3), activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model3.add(MaxPooling2D(pool_size=(2, 2)))

# flatten for fully connected classification layer
model3.add(Flatten())

# --- fully connected layer ---
model3.add(Dense(128, activation='relu',
                  kernel_initializer = init.he_normal(109)))

# --- fully connected layer ---
model3.add(Dense(64, activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# --- classification ---
model3.add(Dense(num_classes, activation='sigmoid'))

# prints out a summary of the model architecture
model3.summary()

```

| Layer (type)                   | Output Shape        | Param # |
|--------------------------------|---------------------|---------|
| conv2d_5 (Conv2D)              | (None, 122, 79, 32) | 4736    |
| max_pooling2d_5 (MaxPooling2D) | (None, 61, 39, 32)  | 0       |
| conv2d_6 (Conv2D)              | (None, 57, 35, 32)  | 25632   |
| max_pooling2d_6 (MaxPooling2D) | (None, 28, 17, 32)  | 0       |
| conv2d_7 (Conv2D)              | (None, 26, 15, 64)  | 18496   |

# CNN\_A2

April 26, 2017

## 1 CS109B Project Group 26 - Deep Learning

## Main Specifications - Ver1

**\*\* Data preprocessing:\*\***

- Channel rearrangement
- 14 labels multi label classification
- 20% validation
- Centered features

**\*\* Main Architecture \*\***

- Multi layer CNN :

- Conv2D - Relu depth 32 and 7x7 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 32 and 5x5 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 16 and 3x3 kernel - He uniform initialization - MaxPool 2x2 - Conv2D - Relu depth 16 and 3x3 kernel - He uniform initialization - FC 128 , Relu , He uniform initialization - FC 64 , Relu , He uniform initialization

- SGD optimizer with 1e-6 learning rate and 0.99 momentum - Binary cross entropy loss function - Batch size 64 - Training convergence after 50 epochs

### 1.0.1 Import Modules

```
In [1]: import requests
import json
import time
import itertools
import wget
import os
import pickle
import numpy as np

import random
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```

import seaborn as sns
from sklearn.cluster.bicluster import SpectralCoclustering
from sklearn.metrics import precision_recall_curve
import scipy

sns.set_style('white')
import tensorflow as tf
import pandas as pd
import keras
from keras.optimizers import SGD
from keras.models import Sequential
from keras.layers import Dense,Dropout,Flatten
from keras.layers import Conv2D, MaxPooling2D
import keras.initializers as init
from keras import backend as K
from keras.models import load_model

```

Using TensorFlow backend.

## 1.0.2 Import the Preprocessed Data from the customized AWS Image EBS Storage

```
In [2]: x_train_dict = pickle.load(open('training_num.pik' , 'rb'))
```

```
In [3]: # Extract the list of input tensors
x_train_raw = x_train_dict['images'][:9988]
```

```

# Transform them into numpy tensor
x_train = np.array(x_train_raw)

```

```

In shape
print x_train.shape

```

```
(9988, 128, 85, 3)
```

### Preprocess the training data

```

In [4]: # Extract the image rows
img_rows = x_train.shape[1]

# Extract the image columns
img_cols = x_train.shape[2]

# Re-arrange according to the configured order of channels
# If Channels first
if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 3, img_rows, img_cols)
    input_shape = (3, img_rows, img_cols)

```

```

# If channels last
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 3)
    input_shape = (img_rows, img_cols, 3)

```

### Normalize the data

```

In [5]: # transform into float
x_train = x_train.astype('float32')

# normalize
x_train /= 255

print 'x_train shape:', x_train.shape
print x_train.shape[0], 'train samples'

```

```

x_train shape: (9988, 128, 85, 3)
9988 train samples

```

### Pre process the label

```

In [6]: # Read From File
y_raw = pd.read_csv('Genres_labels_All_cleaned.csv')

# Select Labels from file
y_train = y_raw.iloc[:, 1:-1].values

# Define the shape of the labels
num_classes = y_train.shape[1]

```

### 1.0.3 Main Architecture Build up

```

In [12]: # create an empty network model
model1 = Sequential()

# --- input layer ---
model1.add(Conv2D(16, kernel_size=(7, 7), activation='relu', input_shape=input_shape,
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model1.add(MaxPooling2D(pool_size=(2, 2)))

# ---- Conv Layer ---
model1.add(Conv2D(16, kernel_size=(5, 5), activation='relu',
                  kernel_initializer = init.he_normal(109)))

```

```

# --- max pool ---
model1.add(MaxPooling2D(pool_size=(2, 2)))

# --- Conv layer ---
model1.add(Conv2D(32, kernel_size=(3, 3), activation='relu',
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model1.add(MaxPooling2D(pool_size=(2, 2)))

# --- Conv layer ---
model1.add(Conv2D(32, kernel_size=(3, 3), activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# --- max pool ---
model1.add(MaxPooling2D(pool_size=(2, 2)))

# flatten for fully connected classification layer
model1.add(Flatten())

# --- fully connected layer ---
model1.add(Dense(128, activation='relu',
                  kernel_initializer = init.he_normal(109)))

# --- fully connected layer ---
model1.add(Dense(64, activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# --- classification ---
model1.add(Dense(num_classes, activation='sigmoid'))

# prints out a summary of the model architecture
model1.summary()

```

| Layer (type)                   | Output Shape        | Param # |
|--------------------------------|---------------------|---------|
| conv2d_8 (Conv2D)              | (None, 122, 79, 16) | 2368    |
| max_pooling2d_8 (MaxPooling2D) | (None, 61, 39, 16)  | 0       |
| conv2d_9 (Conv2D)              | (None, 57, 35, 16)  | 6416    |
| max_pooling2d_9 (MaxPooling2D) | (None, 28, 17, 16)  | 0       |
| conv2d_10 (Conv2D)             | (None, 26, 15, 32)  | 4640    |



# CNN\_BN

April 26, 2017

## 1 CS109B Project Group 26 - Deep Learning

## Main Specifications - Ver4

\*\* Data preprocessing:\*\*

- Channel rearrangement
- 14 labels multi label classification
- data augmentation by shift and zoom for 8000 training samples and 2000 test samples
- Centered features

\*\* Main Architecture \*\*

- Multi layer CNN with Batch Normalization Layers :

- Conv2D - Relu depth 32 and 7x7 kernel - He uniform initialization - Batch Norm - MaxPool  
2x2 - Conv2D - Relu depth 32 and 5x5 kernel - He uniform initialization - Batch Norm - MaxPool  
2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - Batch Norm - MaxPool  
2x2 - Conv2D - Relu depth 64 and 3x3 kernel - He uniform initialization - Batch Norm - FC 128  
, Relu , He uniform initialization - Batch Norm - FC 64 , Relu , He uniform initialization - Batch  
Norm

- SGD optimizer with 1e-4 learning rate and 0.99 momentum - Binary cross entropy loss function - Batch size 128 - Training convergence after 50 epochs

### 1.0.1 Import Modules

```
In [1]: import requests
import json
import time
import itertools
import wget
import os
import pickle
import numpy as np

import random
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```

import seaborn as sns
from sklearn.cluster.bicluster import SpectralCoclustering
from sklearn.metrics import precision_recall_curve
import scipy

sns.set_style('white')
import tensorflow as tf
import pandas as pd
import keras
from keras.optimizers import SGD, Adam
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
import keras.initializers as init
from keras.preprocessing.image import ImageDataGenerator
from keras import backend as K
from keras.models import load_model

```

Using TensorFlow backend.

```
In [2]: x_train_dict = pickle.load(open('training_num.pik' , 'rb'))
```

```
In [3]: train_split = 8000
```

```
In [4]: x_train_raw = x_train_dict['images']
```

```

x_train = np.array(x_train_raw)[:train_split,: , : , :]
x_test  = np.array(x_train_raw)[train_split:,: , : , :]

print x_train.shape

```

```
(8000, 128, 85, 3)
```

```
In [5]: img_rows = x_train.shape[1]
```

```
img_cols = x_train.shape[2]
```

```

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 3, img_rows, img_cols)
    input_shape = (3, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 3)
    input_shape = (img_rows, img_cols, 3)

```

```
In [6]: x_train = x_train.astype('float32')
        x_test  = x_test.astype('float32')

        x_train /= 255.0
        x_test  /= 255.0

        print 'x_train shape:', x_train.shape
        print x_train.shape[0], 'train samples'

        print 'x_test shape:', x_test.shape
        print x_test.shape[0], 'test samples'
```

```
x_train shape: (8000, 128, 85, 3)
8000 train samples
x_test shape: (1988, 128, 85, 3)
1988 test samples
```

```
In [59]: y_raw = pd.read_csv('Genres_labels_All_cleaned.csv')

        y_train = y_raw.iloc[:, 1:-1].values[:train_split, :]
        y_test  = y_raw.iloc[:, 1:-1].values[train_split:, :]

        num_classes = y_train.shape[1]

        print 'number of classes: ', num_classes
```

```
number of classes: 14
```

```
In [60]: datagen = ImageDataGenerator(
        featurewise_center=True,
        featurewise_std_normalization=True,
        width_shift_range=0.2,
        height_shift_range=0.2,
        zoom_range = 0.5,
        fill_mode = 'wrap')

        datagen.fit(x_train)

        datagen.fit(x_test)
```

```
In [74]: # create an empty network model
        modela = Sequential()

        # --- input layer ---
        modela.add(Conv2D(32, kernel_size=(7, 7), activation='relu', input_shape=input_shape
                           kernel_initializer = init.he_normal(109)))
```

```

# -----Batch Normalization -----
modela.add(BatchNormalization())

# --- max pool ---
modela.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# --- Conv Layer ---
modela.add(Conv2D(32, kernel_size=(5, 5), activation='relu',
                  kernel_initializer = init.he_normal(109)))

# -----Batch Normalization -----
modela.add(BatchNormalization())

# --- max pool ---
modela.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# --- Conv layer ---
modela.add(Conv2D(64, kernel_size=(3, 3), activation='relu',
                  kernel_initializer = init.he_normal(109)))

# -----Batch Normalization -----
modela.add(BatchNormalization())

# --- max pool ---
modela.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# --- Conv layer ---
modela.add(Conv2D(64, kernel_size=(3, 3), activation='relu' ,
                  kernel_initializer = init.he_normal(109)))

# -----Batch Normalization -----
modela.add(BatchNormalization())

# --- max pool ---
modela.add(MaxPooling2D(pool_size=(2, 2)))

# -----
# flatten for fully connected classification layer
modela.add(Flatten())

# --- fully connected layer ---

```

```

modela.add(Dense(128, activation='relu',
                 kernel_initializer = init.he_normal(109)))

# -----Batch Normalization -----
modela.add(BatchNormalization())

# -----

# --- fully connected layer ---
modela.add(Dense(64, activation='relu' ,
                 kernel_initializer = init.he_normal(109)))

# -----Batch Normalization -----
modela.add(BatchNormalization())

# -----

# --- classification ---
modela.add(Dense(num_classes, activation='sigmoid'))

# prints out a summary of the model architecture
modela.summary()

```

| Layer (type)                                 | Output Shape        | Param # |
|--|---------------------|---------|
| conv2d_48 (Conv2D)                           | (None, 122, 79, 32) | 4736    |
| batch_normalization_63 (Batch Normalization) | (None, 122, 79, 32) | 128     |
| max_pooling2d_48 (MaxPooling2D)              | (None, 61, 39, 32)  | 0       |
| conv2d_49 (Conv2D)                           | (None, 57, 35, 32)  | 25632   |
| batch_normalization_64 (Batch Normalization) | (None, 57, 35, 32)  | 128     |
| max_pooling2d_49 (MaxPooling2D)              | (None, 28, 17, 32)  | 0       |
| conv2d_50 (Conv2D)                           | (None, 26, 15, 64)  | 18496   |
| batch_normalization_65 (Batch Normalization) | (None, 26, 15, 64)  | 256     |
| max_pooling2d_50 (MaxPooling2D)              | (None, 13, 7, 64)   | 0       |
| conv2d_51 (Conv2D)                           | (None, 11, 5, 64)   | 36928   |
| batch_normalization_66 (Batch Normalization) | (None, 11, 5, 64)   | 256     |

```
In [89]: from __future__ import print_function

import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation
from keras.optimizers import SGD

import matplotlib
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('white')

import numpy as np
```

First we load in the data. There are two unique datasets in this analysis because one includes only the RGB color data and the other utilizes all of the quantitative metadata we scraped.

```
In [98]: # getting RGB data for MLP

x_train = np.genfromtxt('train1_RGB.csv', delimiter=',', skip_header = 1)
y_train = np.genfromtxt('train1_y.csv', delimiter=',', skip_header = 1)

x_train_full = np.genfromtxt('train1_x.csv', delimiter=',', skip_header = 1)
# print(x_train_full)

x_test = np.genfromtxt('test1_RGB.csv', delimiter=',', skip_header = 1)
y_test = np.genfromtxt('test1_y.csv', delimiter=',', skip_header = 1)

x_test2 = np.genfromtxt('test2_x.csv', delimiter=',', skip_header = 1)
y_test2 = np.genfromtxt('test2_y.csv', delimiter=',', skip_header = 1)
```

In **model** we only utilize the color data (RGB means and sd for posters). Here, we utilized the Sequential model in Keras in order to create an MLP model for multi-level softmax classification. Instead of using the "relu" activation parameter, we chose to utilize the "sigmoid" activation because it was better across the board. We experimented with the learning rate by hand to examine the process along the epochs. Experimenting with (1e-5, 1e-4, 1e-3, .01, .1) we concluded that .01 was the best learning rate. We decided on the .01 rate by looking at the training data even though when evaluated with the test data the total accuracy was roughly the same. We also experimented with some momentum values (.5 - .99), but .9 seemed to be the best. They were often roughly the same, but other values seemed to have more volatile output accuracies. So, .9 seemed to be a pretty typical value that was reliable.

Overall, the accuracy for the color vectors is not very good. On the test set it comes out to about 20% accuracy. However, this is reasonable because our color analysis is extremely elementary. We scraped the mean value for each of the three RGB vectors and also recorded its standard deviation. You can imagine many scenarios where this kind of analysis fails to pick up differences between genres or the character of the movie, but it may have some predictive power when added to a larger deep learning model for the posters themselves.

```
In [91]: model = Sequential()

# try relu or sigmoid
model.add(Dense(64, activation = 'sigmoid', input_dim = 6))
model.add(Dropout(.5))
model.add(Dense(64, activation = 'sigmoid'))
model.add(Dropout(.5))
model.add(Dense(18, activation = 'softmax'))
```

```
In [131]: learn = .01
decay_rate = 1e-6
mom = .9
sgd = SGD(lr = learn, decay = decay_rate, momentum = mom , nesterov = True)
model.compile(loss = 'categorical_crossentropy', optimizer = sgd, metrics = ['accuracy'])
```

```
In [ ]: import requests
import json
import time
import itertools
import wget
import os
import pickle
import numpy as np

import random
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

import seaborn as sns
from sklearn.cluster.bicluster import SpectralCoclustering
from sklearn.metrics import precision_recall_curve
import scipy

sns.set_style('white')
import tensorflow as tf
import pandas as pd
import keras
from keras.optimizers import SGD, Adam
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
import keras.initializers as init
from keras.preprocessing.image import ImageDataGenerator
from keras import backend as K
from keras.models import load_model
```

```
In [ ]: x_train_dict = pickle.load(open('training_num.pik' , 'rb'))
```

```
In [ ]: x_train_raw = x_train_dict['images']

x_train = np.array(x_train_raw)

print x_train.shape
```

```
In [ ]: img_rows = x_train.shape[1]

img_cols = x_train.shape[2]

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 3, img_rows, img_cols)
    input_shape = (3, img_rows, img_cols)

else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 3)
    input_shape = (img_rows, img_cols, 3)
```



```
In [ ]: x_train = x_train.astype('float32')

x_train /= 255

print 'x_train shape:', x_train.shape
print x_train.shape[0], 'train samples'
```

```
In [ ]: y_raw = pd.read_csv('Genres_labels_All_cleaned.csv')

y_train = y_raw.iloc[:, 1:-1].values

num_classes = y_train.shape[1]
```

```
In [ ]: datagen = ImageDataGenerator(
    featurewise_center=True,
    featurewise_std_normalization=True,
    width_shift_range=0.2,
    height_shift_range=0.2,
    zoom_range = 0.5,
    fill_mode = 'wrap')

datagen.fit(x_train)
```

```
In [ ]: # create an empty network model
model = Sequential()

model.add(Dense(64, activation='relu', input_shape=input_shape))
# this is our hidden layer

model.add(Dense(64, activation='relu'))
# and an output layer

model.add(Dense(8, activation='softmax'))

# prints out a summary of the model architecture
model.summary()
```

```
In [ ]: ada = Adam(lr=0.01)
model.compile(loss='binary_crossentropy',
              optimizer=ada,
              metrics=['accuracy'])
```

```
In [ ]: batch_size = 64
epochs = 20
```

```
In [ ]: history = model.fit_generator(datagen.flow(x_train, y_train,
                                                    batch_size=batch_size),
                                     steps_per_epoch=len(x_train) / batch_size,
                                     epochs=epochs)
```

```
In [ ]: plt.plot(history.history['acc'])  
plt.xlabel("epoch")  
plt.ylabel("accuracy")
```

```
In [ ]: import h5py as h5py
```

```
In [ ]: model.save('mlp_var1.h5')
```

```
In [ ]: Acc_mlp_var1 = pd.DataFrame(history.history['acc'] , columns = ['Accurac  
y'])
```

```
In [ ]: Acc_mlp_var1.to_csv('mlp_v1.csv')
```